

```
In [ ]: %load_ext autoreload
        %autoreload 2
```

## Imports

```
In [ ]: import pathlib
import random
import time

import matplotlib.pyplot as plt
import tqdm
import torch
from torch import nn, optim
from torch.nn.functional import interpolate
from torchvision.utils import save_image
from torch.utils.data import DataLoader
from torchvision.transforms.functional import adjust_brightness

from super_resolution.src.sen2venus_dataset import (
    create_train_test_split,
)
from super_resolution.src.visualization import plot_gallery
from super_resolution.src.srgan import SRResNet, Discriminator
```

```
In [ ]: DATA_DIR = pathlib.Path("C:/Users/Mitch/stat3007_data")
SITES_DIR = DATA_DIR / "sites"
PREPROCESSING_DIR = DATA_DIR / "preprocessing"
RESULTS_DIR = DATA_DIR / "results"
```

## Prepare Data

```
In [ ]: sites = {
    "SO2",
    "FR-BIL",
    "NARYN",
}
train_patches, test_patches = create_train_test_split(
    str(SITES_DIR) + "\\ ", sites=sites
)
print(f"Num train {len(train_patches)}\n" f"Num test {len(test_patches)}")
```

```
Num train 8159
Num test 3498
```

```
In [ ]: def image_transform(x, y):
    x = x[:3, :, :]
    y = y[:3, :, :]

    x = torch.clamp(x, 0, 1)
    y = torch.clamp(y, 0, 1)

    return x, y
```

```
In [ ]: train_patches.set_transform(image_transform)
test_patches.set_transform(image_transform)
```

```
In [ ]: train_loader = DataLoader(train_patches, batch_size=100)
```

```
In [ ]: (low_res, high_res) = next(train_loader.__iter__())
```

```
In [ ]: index = random.randint(0, len(low_res) - 1)
low_res_example = low_res[index]

high_res_example = high_res[index]

plot_gallery(
    [
        adjust_brightness(low_res_example, 2).permute(1, 2, 0),
        adjust_brightness(high_res_example, 2).permute(1, 2, 0),
    ],
    titles=["low res", "high res"],
    xscale=5,
    yscale=5,
)
```



## Training

```
In [ ]: # For clearing GPU memory
1 / 0
```

```
In [ ]: # For clearing GPU memory
import gc

gc.collect()
torch.cuda.empty_cache()
```

```
In [ ]: BATCH_SIZE = 16
GEN_LEARNING_RATE = 0.0001
DISCRIM_LEARNING_RATE = 0.0001
```

```
In [ ]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
torch.cuda.empty_cache()
should_pin_memory = {"cuda": True, "cpu": False}

train_loader = DataLoader(
    train_patches,
    shuffle=True,
    batch_size=BATCH_SIZE,
    drop_last=True,
    pin_memory=should_pin_memory[device.type],
    num_workers=0,
)
```

```
In [ ]: device
```

```
Out[ ]: device(type='cuda')
```

```
In [ ]: generator = SRResNet(scaling_factor=2, n_blocks=8)
```

```
In [ ]: discriminator = Discriminator(n_blocks=2, fc_size=128)
```

```
In [ ]: generator = generator.to(device)
generator = generator.train()
```

```
In [ ]: discriminator = discriminator.to(device)
discriminator = discriminator.train()
```

```
In [ ]: gen_optimizer = optim.Adam(generator.parameters(), lr=GEN_LEARNING_RATE)
discrim_optimizer = optim.Adam(discriminator.parameters(), lr=DISCRIM_LEARNING_RATE)
```

```
In [ ]: # save_file = RESULTS_DIR / "greatsrgan_withoutbatchnorm_firstbeatbicubic3_final.pkl"
# Loaded_experiment = torch.load(save_file, map_location=device)
# generator2 = SRResNet(scaling_factor=2, n_blocks=8).to(device)
# generator2.load_state_dict(Loaded_experiment["gen_state"])
# discriminator.load_state_dict(Loaded_experiment["discrim_state"])
# gen_optimizer.load_state_dict(Loaded_experiment["gen_optimizer_state"])
# discrim_optimizer.load_state_dict(Loaded_experiment["discrim_optimizer_state"])
```

```

# gen_losses = loaded_experiment["gen_losses"]
# discrim_losses = loaded_experiment["discrim_losses"]

In [ ]: pixel_criterion = nn.MSELoss()
        discrim_criterion = nn.BCEWithLogitsLoss()

In [ ]: NUM_EPOCHS = 2000

In [ ]: train_time = 0.0
        gen_losses = []
        discrim_losses = []

In [ ]: for epoch in range(NUM_EPOCHS):
        progress_bar = tqdm.tqdm(train_loader, total=len(train_loader), ncols=100)
        gen_epoch_loss = 0.0
        discrim_epoch_loss = 0.0
        num_batches = 0
        for low_res_batch, high_res_batch in progress_bar:
            num_batches += 1
            start_time = time.time()

            # Push to GPU
            low_res_batch = low_res_batch.to(device)
            high_res_batch = high_res_batch.to(device)

            # Update generator
            #####
            super_resolved = generator(low_res_batch).clamp(0, 1)
            natural_probs = discriminator(super_resolved)

            pixel_loss = pixel_criterion(super_resolved, high_res_batch)
            discrim_loss = discrim_criterion(natural_probs, torch.ones_like(natural_probs))
            gen_loss = pixel_loss + 1e-3 * discrim_loss

            gen_optimizer.zero_grad()
            gen_loss.backward()
            gen_optimizer.step()

            # Update discriminator
            #####
            true_natural_probs = discriminator(high_res_batch)

            # Detach to skip generator computations
            fake_natural_probs = discriminator(super_resolved.detach())

            true_natural_loss = discrim_criterion(
                true_natural_probs, torch.ones_like(true_natural_probs)
            )
            fake_natural_loss = discrim_criterion(
                fake_natural_probs, torch.ones_like(fake_natural_probs)
            )
            adversarial_loss = true_natural_loss + fake_natural_loss

            # Do not let discrim get too much advantage
            if adversarial_loss.item() > 1e-04:
                discrim_optimizer.zero_grad()
                adversarial_loss.backward()
                discrim_optimizer.step()

            # Collect data
            gen_epoch_loss += gen_loss.item()
            discrim_epoch_loss += discrim_loss.item()
            progress_bar.set_postfix(
                epoch=epoch,
                gen_loss=f"{gen_epoch_loss/num_batches:.8f}",
                discrim_loss=f"{discrim_epoch_loss/num_batches:.8f}",
            )

            end_time = time.time()
            train_time += end_time - start_time

        gen_epoch_loss /= len(train_loader)
        discrim_epoch_loss /= len(train_loader)
        gen_losses.append(gen_epoch_loss)
        discrim_losses.append(discrim_epoch_loss)
        print(
            f"Epoch: {epoch} / gen_loss: {gen_epoch_loss:.8f} / discrim_loss: {discrim_epoch_loss:.8f}"
        )

In [ ]: train_time / 60

Out[ ]: 343.2222263852755

In [ ]: len(gen_losses)


```

Out[ ]: 105


```
In [ ]: gen_losses[-5:]
```

```
Out[ ]: [2.8041264508891584e-05,  
        2.8292864500128706e-05,  
        2.7803810764296135e-05,  
        2.7878230204005235e-05,  
        2.7838350199665558e-05]
```

```
In [ ]: plt.plot(gen_losses)  
plt.title("Generator losses")  
plt.xlabel("Epoch")  
plt.ylabel("Perceptual Loss")  
plt.show()
```



```
In [ ]: plt.plot(discrim_losses)  
plt.title("Discriminator losses")  
plt.xlabel("Epoch")  
plt.ylabel("Crossentropy Loss")  
plt.show()
```



## Saving

```
In [ ]: experiment = {  
    "gen_losses": gen_losses,  
    "discrim_losses": discrim_losses,  
    "gen_state": generator.state_dict(),  
    "discrim_state": discriminator.state_dict(),  
}
```

```

"gen_optimizer_state": gen_optimizer.state_dict(),
"discrim_optimizer_state": discrim_optimizer.state_dict(),
"gen_learning_rate": GEN_LEARNING_RATE,
"discrim_learning_rate": DISCRIM_LEARNING_RATE,
}

```

```
In [ ]: save_file = RESULTS_DIR / "greatsrgan_withoutbatchnorm_firstbeatbicubic3_final.pkl"
```

```
In [ ]: if not save_file.exists():
        torch.save(experiment, save_file)
        print(f"Saved to {save_file}")
```

Saved to C:\Users\Mitch\stat3007\_data\results\greatsrgan\_withoutbatchnorm\_firstbeatbicubic3\_final.pkl

## Generated Examples

```
In [ ]: generator = generator.to("cpu")
```

```
In [ ]: loader = DataLoader(train_patches, batch_size=100)
        (low_res, high_res) = next(loader.__iter__())
```

```
In [ ]: import random

index = random.randint(0, len(low_res) - 1)

low_res_example = low_res[index]
high_res_example = high_res[index]
out = generator(low_res_example.unsqueeze(0)).detach().clamp(0, 1)
bicubic_out = interpolate(
    low_res_example.unsqueeze(0),
    size=(256, 256),
    mode="bicubic",
).float()
plot_gallery(
    [
        adjust_brightness(low_res_example, 2).permute(1, 2, 0),
        adjust_brightness(high_res_example, 2).permute(1, 2, 0),
        adjust_brightness(out[0], 2).permute(1, 2, 0),
        adjust_brightness(bicubic_out[0], 2).permute(1, 2, 0),
    ],
    titles=["low res", "high res", "SRGAN", "Bicubic"],
    xscale=5,
    yscale=5,
)
```



```
In [ ]: loader = DataLoader(test_patches, batch_size=300)
        (low_res, high_res) = next(loader.__iter__())
```

```
In [ ]: import random

index = random.randint(0, len(low_res) - 1)

low_res_example = low_res[index]
high_res_example = high_res[index]

out = generator(low_res_example.unsqueeze(0)).detach().clamp(0, 1)[0]
out = adjust_brightness(out, 2)

bicubic_out = (
    interpolate(low_res_example.unsqueeze(0), size=(256, 256), mode="bicubic")
    .float()
    .clamp(0, 1)
)[0]
bicubic_out = adjust_brightness(bicubic_out, 2)

low_res_example = adjust_brightness(low_res_example, 2)
high_res_example = adjust_brightness(high_res_example, 2)

plot_gallery(
```

```
[
    low_res_example.permute(1, 2, 0),
    high_res_example.permute(1, 2, 0),
    out.permute(1, 2, 0),
    bicubic_out.permute(1, 2, 0),
],
titles=["low res", "high res", "SRGAN", "Bicubic"],
xscale=5,
yscale=5,
)
```



## Save images

```
In [ ]: IMAGE_DIR = RESULTS_DIR / "images"
```

```
In [ ]: index
```

```
Out[ ]: 250
```

```
In [ ]: srgan_file = IMAGE_DIR / f"{index}srgan.png"
bicubic_file = IMAGE_DIR / f"{index}bicubic.png"
low_res_file = IMAGE_DIR / f"{index}lowres.png"
high_res_file = IMAGE_DIR / f"{index}highres.png"
```

```
In [ ]: save_image(low_res_example, low_res_file)
save_image(high_res_example, high_res_file)
save_image(out, srgan_file)
save_image(bicubic_out, bicubic_file)
```

## Metrics

```
In [ ]: from super_resolution.src.testing import compute_metrics
```

```
In [ ]: # 256 is largest I can handle on cpu
metric_loader = DataLoader(test_patches, batch_size=min(len(test_patches), 256))
```

```
In [ ]: metrics = compute_metrics(lambda x: generator(x).clamp(0, 1), metric_loader)
```

```
100%|████████████████████████████████████████| 14/14 [15:47<00:00, 67.71s/it]
```

```
In [ ]: metrics
```

```
Out[ ]: Metrics(mse=3.124712371313113e-05, psnr=45.06763458251953, ssim=0.9879355771200997, fid=0.9709701538085938)
```

```
In [ ]: bicubic_metrics = compute_metrics(
    lambda x: interpolate(x, size=(256, 256), mode="bicubic"), metric_loader
)
```

```
In [ ]: bicubic_metrics
```

```
Out[ ]: Metrics(mse=4.024813668885534e-05, psnr=43.96280016217913, ssim=0.9836511441639492, fid=tensor(1.3964))
```

```
In [ ]:
```